Akka Streams: a match in heaven for Reactive Systems

Henrik Engström, Senior Software Engineer

UJUG presentation 21/09/2017

AGENDA - UJUG 21/9/17

- Traditional systems architecture
- Reactive systems architecture
- Akka Actors
- Reactive Streams
- Akka Streams
- Alpakka
- Actors and Streams

PDF/Code found here: **github.com/henrikengstrom/ujug2017**

Traditional systems architecture

Traditional systems architecture

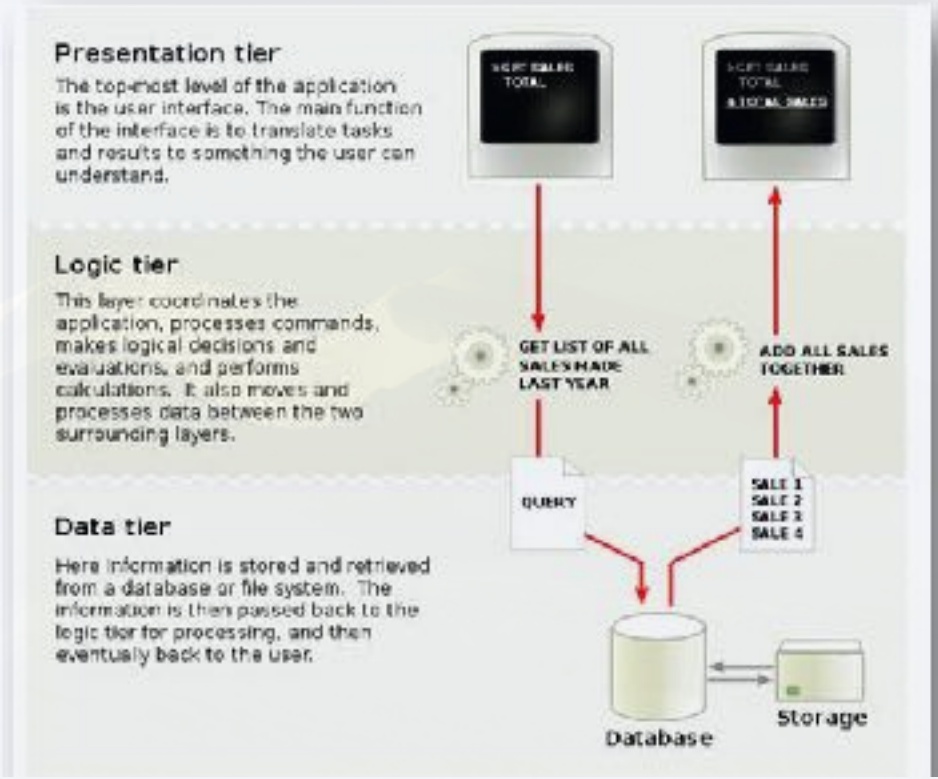

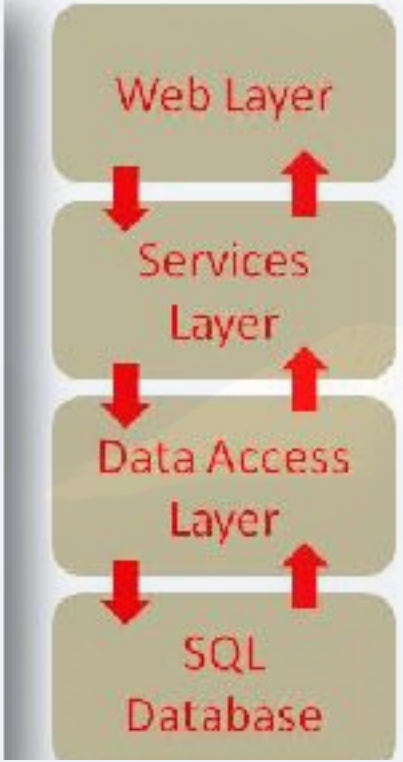

TSA - frameworks and servers

GlassFish

spring

Apache
SOFTWARE FOUNDATION

WebSphere
Application Server

HIBERNATE

TSA - cloud

kubernetes

TSA - some struggles

- Hard to engineer systems to:
  - withstand load
  - always stay up
  - be performant
  - utilize HW to a maximum

- What can we do about it?

Reactive systems

The Reactive Manifesto

- http://www.reactivemanifesto.org/
- September 16, 2014
- +20k signatures
- Four traits:
  - *Responsive*
  - *Resilient*
  - *Elastic*
  - *Message driven*

Responsive

Resilient

Elastic

Message driven

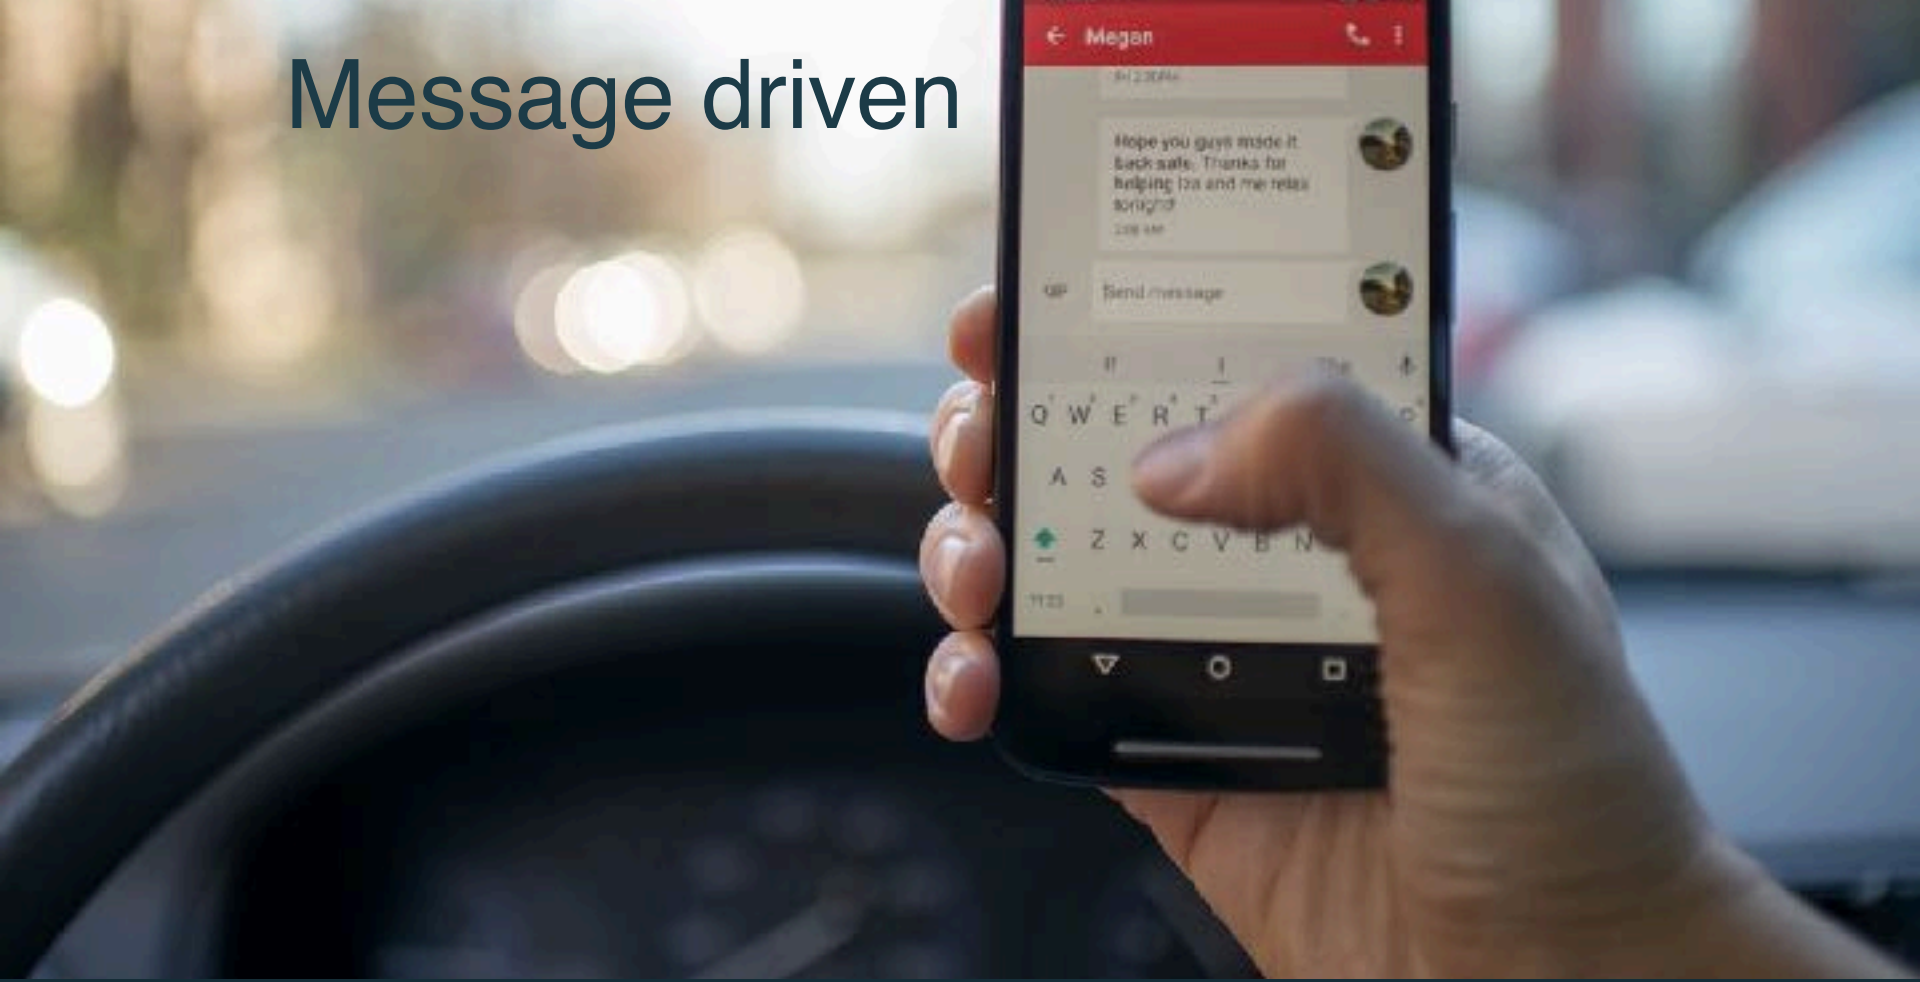

akka Actors

- Actor
  - Mailbox
  - Behavior
  - State
- ActorSystem
  - Thread pools
  - Configuration
  - System actors
- Messages
- JVM

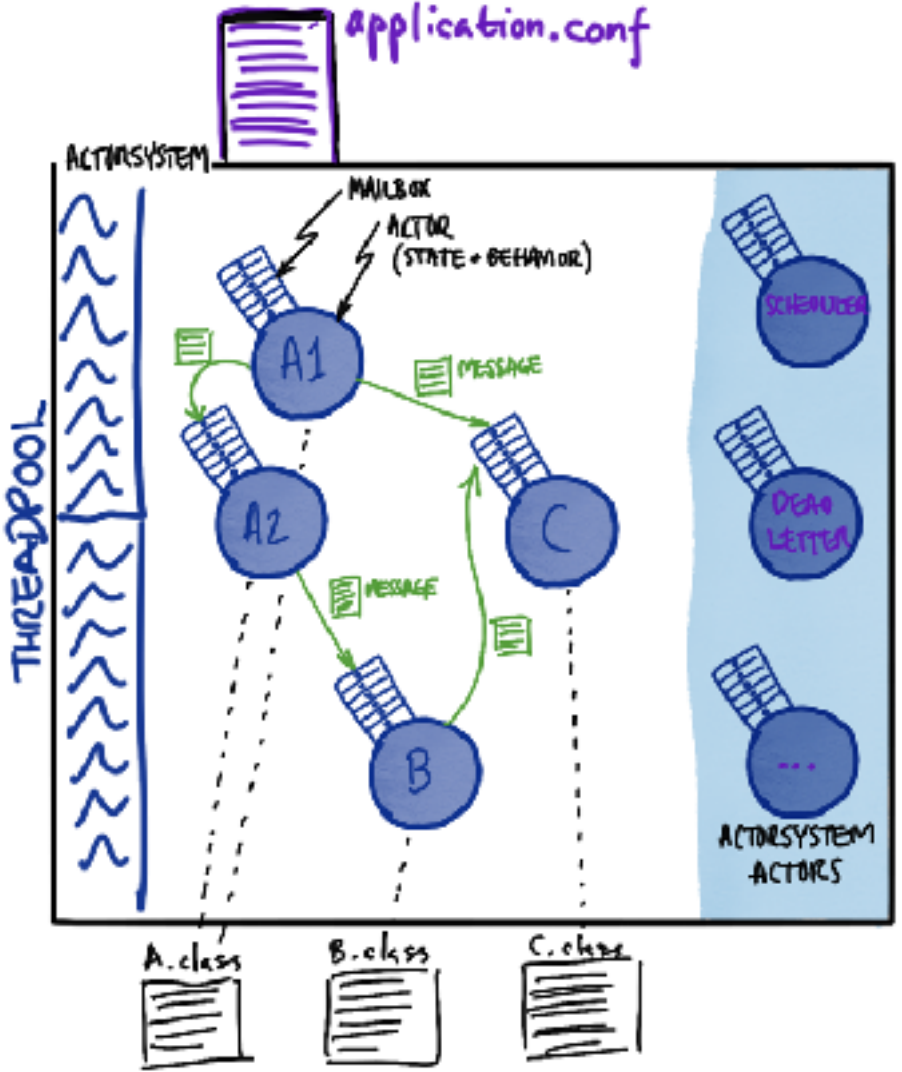

Defining actors

```
import akka.actor.AbstractActor;
import akka.event.Logging;
import akka.event.LoggingAdapter;
public class MyActor extends AbstractActor {
    private final LoggingAdapter log = Logging.getLogger(getContext().getSystem(), this);
    private int msgs = 0;
     @Override
      public Receive createReceive() {
          return receiveBuilder()
           .match(String.class, s -> {
               msgs++;
               log.info("Received message: {}, msg number: {}", s, msgs);
           }).build()
       }
}
```


ActorSystem

```
final ActorSystem actorSystem =
    ActorSystem.create("AS");
```

Creating actors

```
final ActorRef myActor =
 actorSystem.actorOf(Props.create(MyActor.class));

final ActorRef myActor =
  actorSystem.actorOf(Props.create(MyActor.class),
        "myActor");

final ActorRef myActor =
    actorSystem.actorOf(MyActor.props(),"myActor");
```

Defining messages

```
package example;
import akka.actor.AbstractActor;
public class MyActor extends AbstractActor {
    // implement createReceive

     public static class SomeMessage {
         public String someValue;
         public SomeMessage(String someValue) {
             this.someValue = someValue;
         }
        // implement hashCode, equals, toString
     }
}
```

Sending messages

```
final ActorRef myActorRef =
    actorSystem.actorOf(MyActor.props(), "myActor");
myActorRef.tell(
    new MyActor.SomeMessage("something"),
ActorRef.noSender());

// Or if in the context of an actor
final ActorRef myActorRef =
    getContext().actorOf(MyActor.props(), "myActor");
myActorRef.tell(
    new MyActor.SomeMessage("something"), getSender());
```

Other important concepts for this demo

```
// the "ask" pattern
CompletionStage<Object> futureResult =
    ask(myActorRef,
        new MyActor.SomeMessage("something"), 1000);
// configuration - src/main/resources/application.conf
// someContext { a-b-c = 123 }
// HOCON - Typesafe Config
int someContextABC =
   getContext()
        .getSystem()
            .settings()
                .config()
                    .getInt("someContext.a-b-c");
```

Example app: Microservices (of course…)

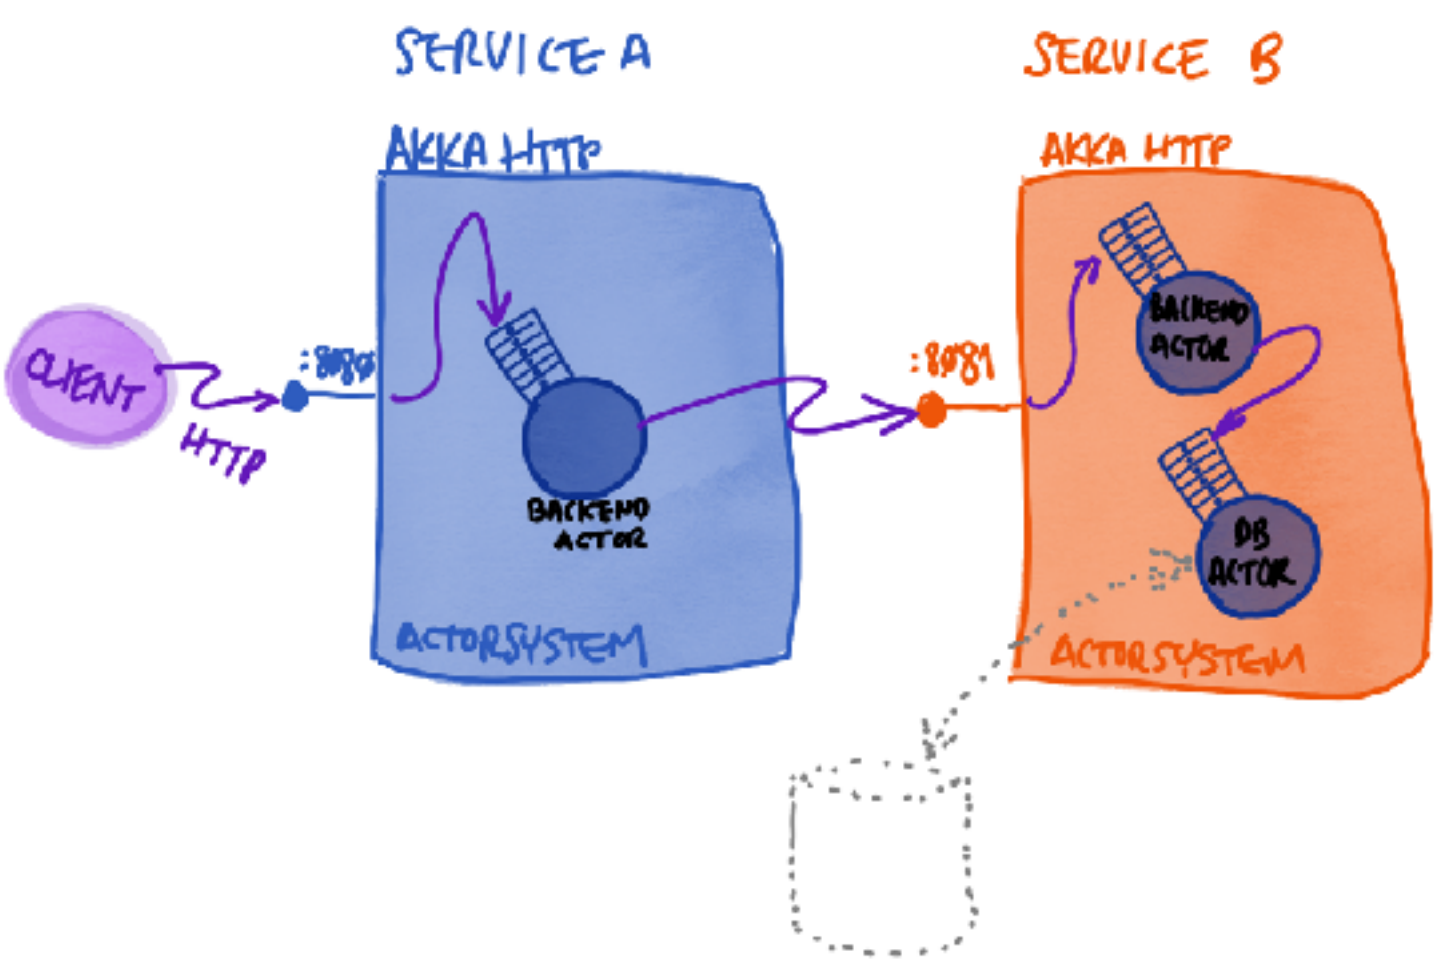
SERVICE A
AKKA HTTP
CLIENT
HTTP
:8080
BACKEND ACTOR
ACTORSYSTEM
SERVICE B
AKKA HTTP
:8081
BACKEND ACTOR
DB ACTOR
ACTORSYSTEM

Example app - coding time!

Reactive Streams

Reactive Streams - reactivestreams.org

"Reactive Streams is an initiative to provide a standard for ***asynchronous stream processing*** with ***non-blocking back pressure***. This encompasses efforts aimed at runtime environments (JVM and JavaScript) as well as network protocols."

What is a "stream"?

- A possibly infinite set of datum
- Processed element by element
  - Could be *Byte* but more useful *<T>*
- Asynchronous processing
  - Sender and receiver are decoupled
  - Asynchronous boundaries (between *threads*)
  - Network boundaries (between *nodes*)

Reactive Streams

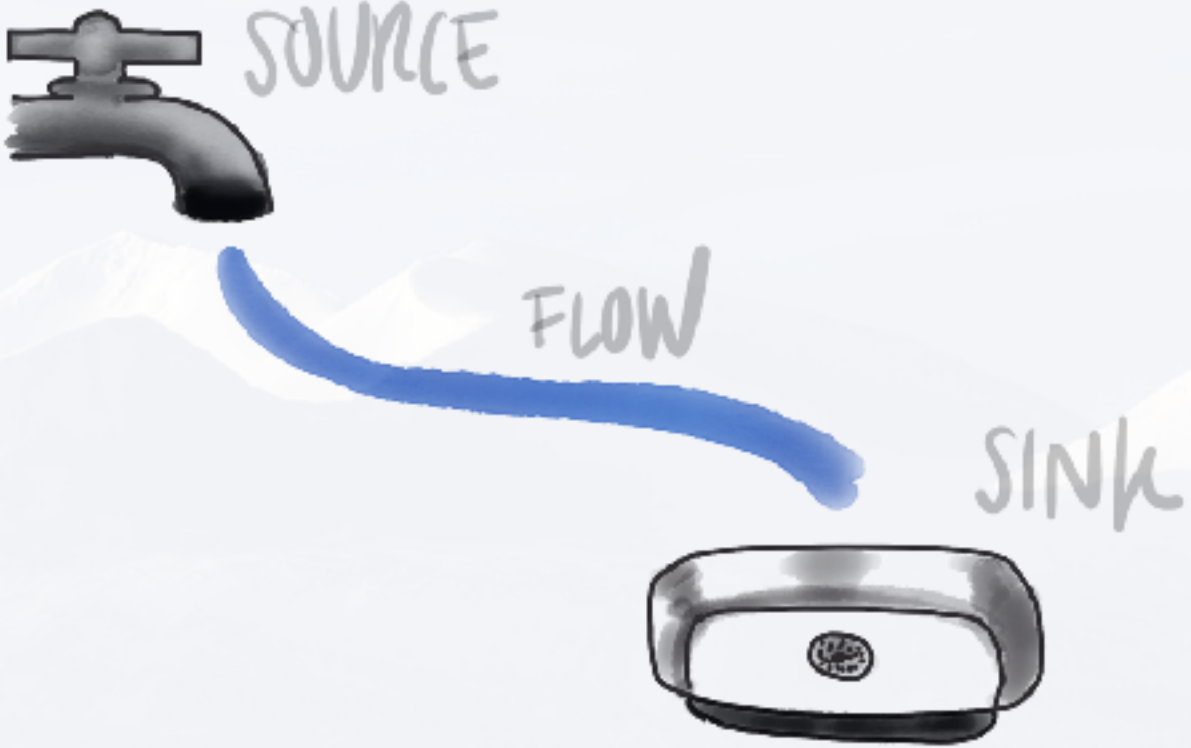

RS: Alternative representation

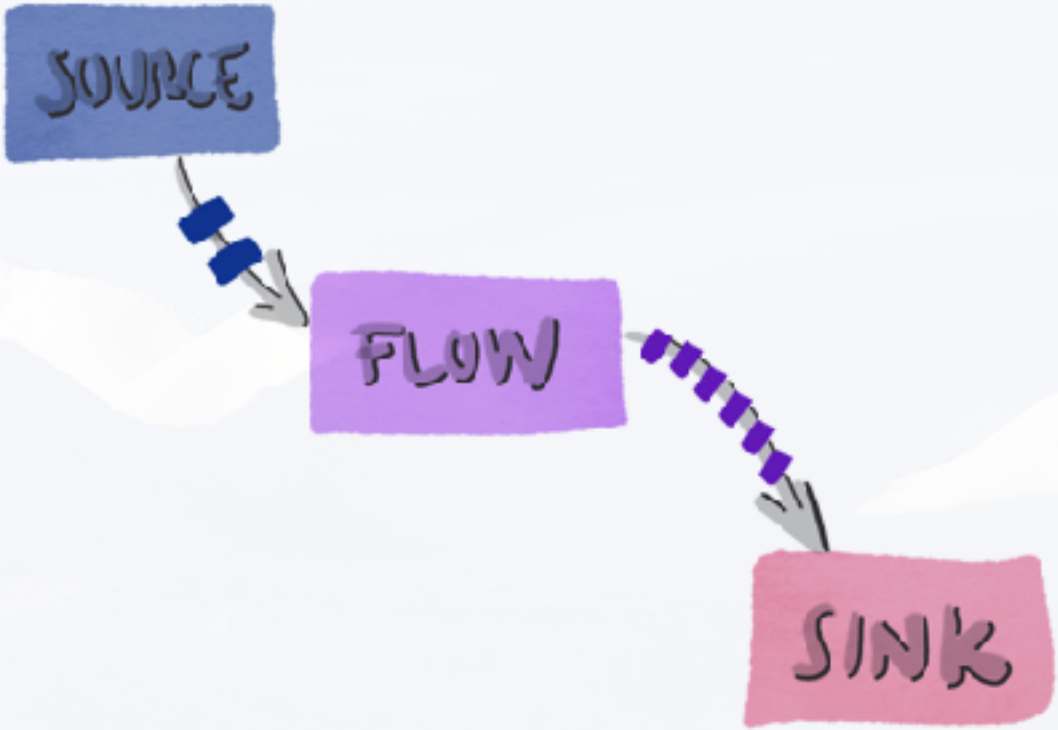

Asynchronous

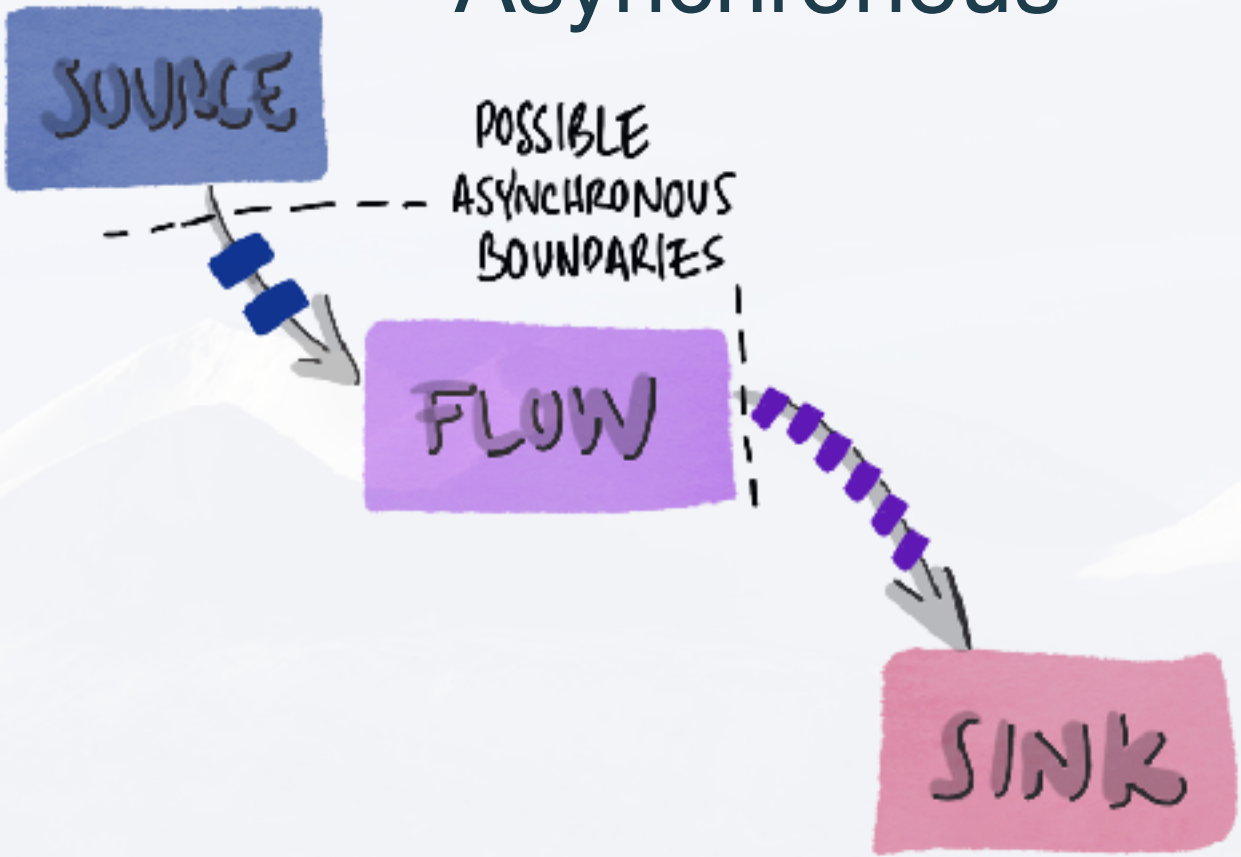

Fast Source

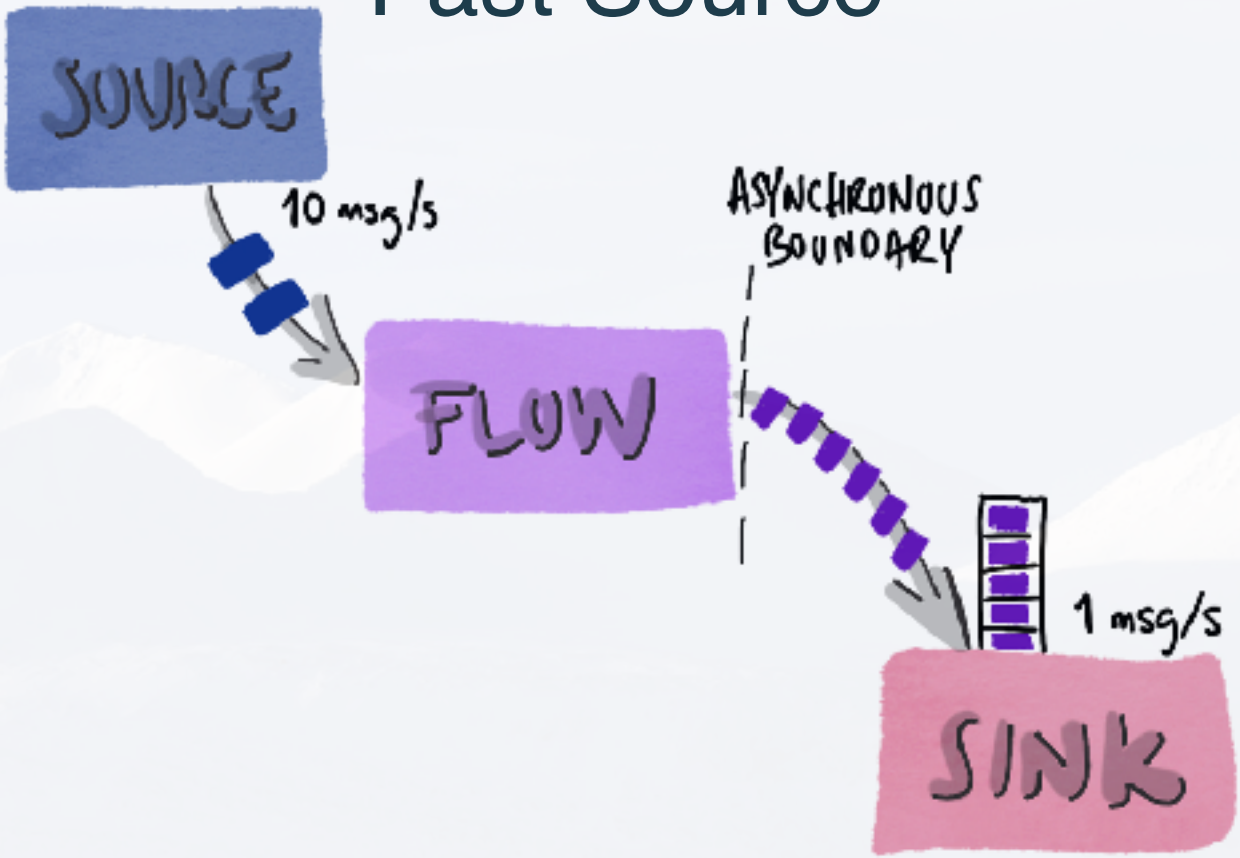

OOM

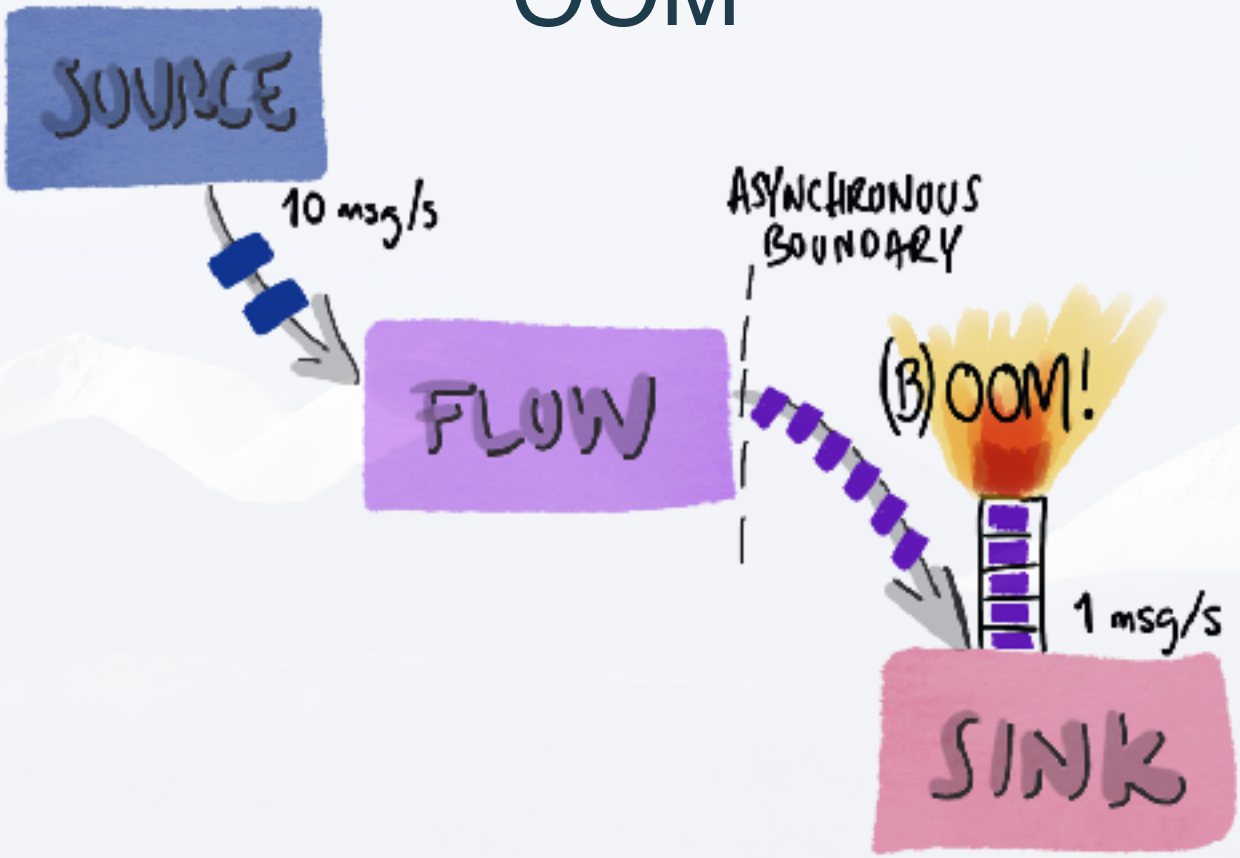

Lost msgs

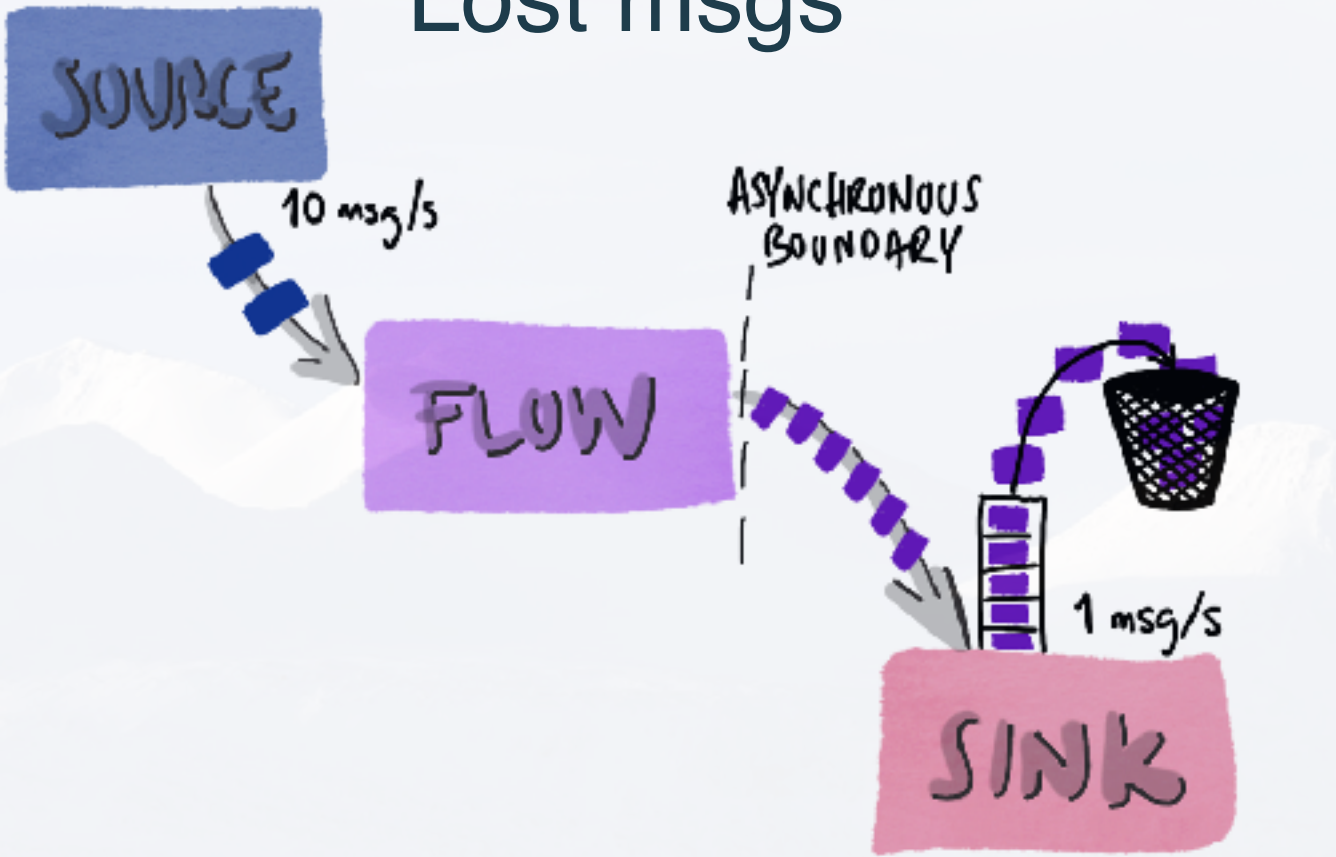

Back pressure

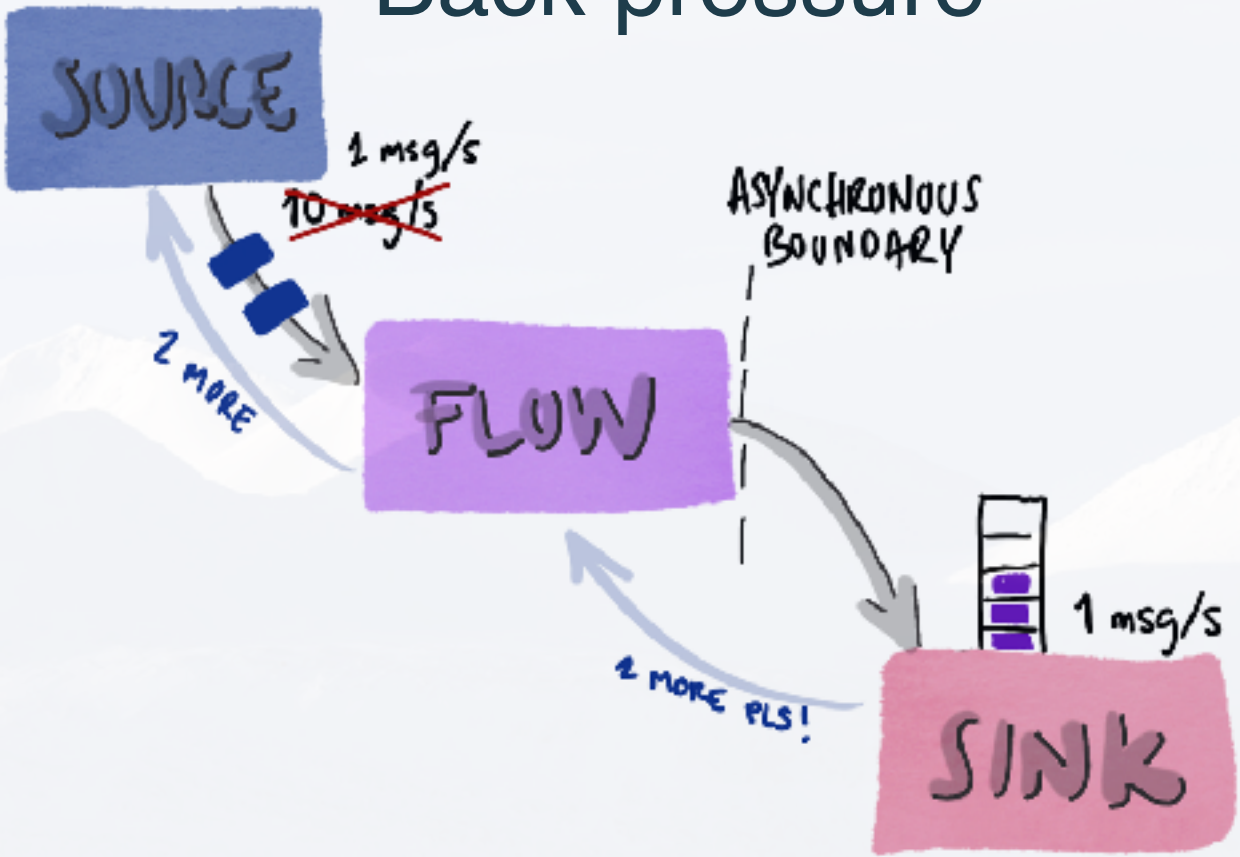

Reactive streams specification

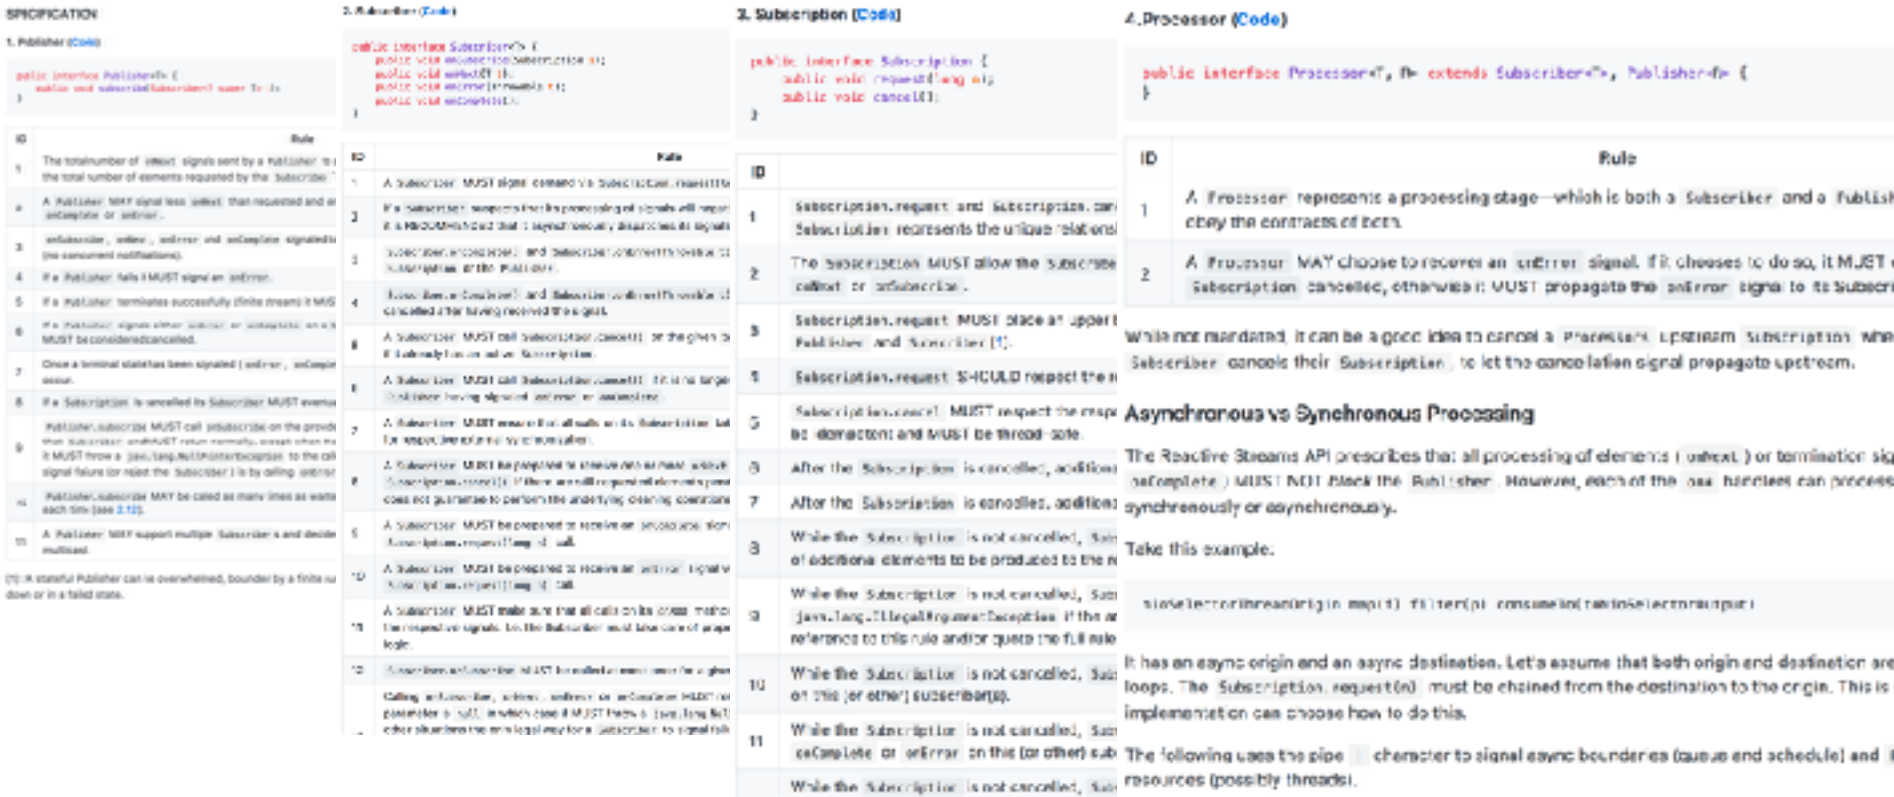

Reactive Streams interoperability

- Vert.x
- RxJava
- Reactor
- Akka Streams

RxJava <-> Vert.x <-> Akka Streams

JDK9 - *java.util.concurrent.Flow*

- Flow.Publisher<T> : *Source*
- Flow.Processor<T, R> : *Flow*
- Flow.Subscriber<T> : *Sink*

Akka Streams

Akka Streams in 60s

```
final ActorSystem actorSystem = ActorSystem.create();
final Materializer materializer =
    ActorMaterializer.create(actorSystem);
final Source<Integer, NotUsed> source =
    Source.range(0, 10);
final Flow<Integer, String, NotUsed> flow =
    Flow.fromFunction((Integer i) -> i.toString());
final Sink<String, CompletionStage<Done>> sink =
    Sink.foreach(s -> System.out.println("Number: " + s));
final RunnableGraph runnable =
    source.via(flow).to(sink);
runnable.run(materializer);
```

Akka Streams in 60s

```
implicit val actorSystem = ActorSystem()
implicit val materializer = ActorMaterializer()

val source = Source(0 to 10)
val flow = Flow[Int].map(_.toString)
val sink =
  Sink.foreach[String](s => println(s"Number: $s"))
val runnable = source.via(flow).to(sink)
runnable.run()
```

Akka Streams in **20**s

```
final ActorSystem actorSystem = ActorSystem.create();
final Materializer materializer =
    ActorMaterializer.create(actorSystem);

Source.range(0, 10)
        .map(Object::toString)
        .runForeach(s ->
            System.out.println("Number: " + s),
                materializer);
```

Graph stages

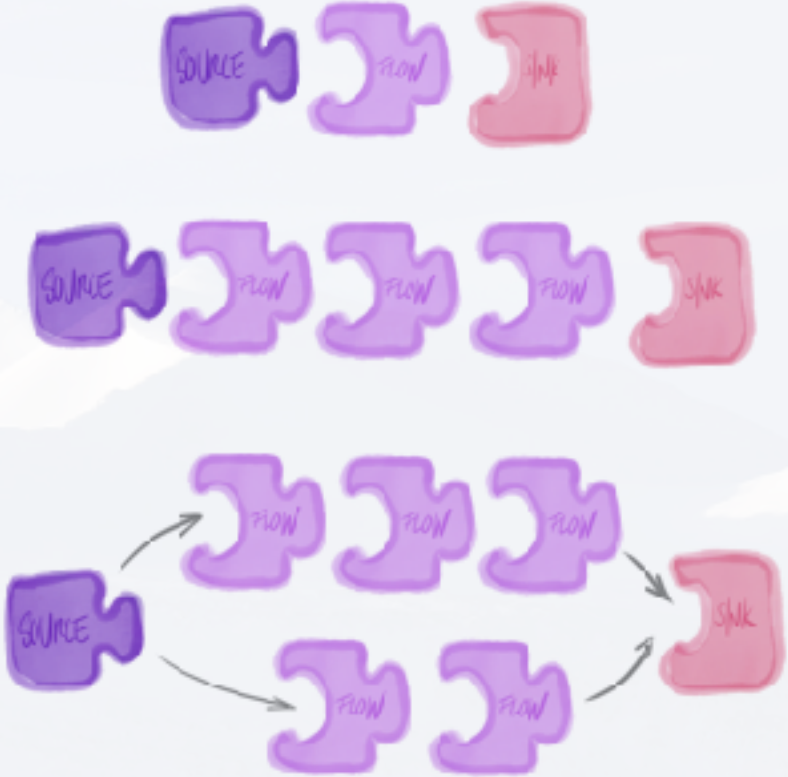

Materialization

Blueprint
(Graph)

Internal representation
(Actors)

Creating akka.stream.javadsl.Source

```
static <T> Source<T,ActorRef> actorRef(int bufferSize,
      OverflowStrategy overflowStrategy);
static <O> Source<O,NotUsed> empty();
static <O> Source<O,NotUsed> from(java.lang.Iterable<O> iterable);
static <O> Source<O,NotUsed> fromFuture(scala.concurrent.Future<O> future);
static <O> Source<O,NotUsed> fromIterator(Creator<java.util.Iterator<O>> f);
static Source<java.lang.Integer,NotUsed> range(int start, int end);
static <T> Source<T,NotUsed> single(T element);
// etc.
```

Using akka.streams.javadsl.Flow

```
static <T> Flow<T,T,NotUsed> create();
Flow<In,Out,Mat> drop(long n);
Flow<In,Out,Mat> dropWhile(Predicate<Out> p);
Flow<In,Out,Mat> filter(Predicate<Out> p);
Flow<In,Out,Mat> filterNot(Predicate<Out> p);
<T> Flow<In,T,Mat> fold(T zero, Function2<T,Out,T> f);
<T> Flow<In,T,Mat> foldAsync(T zero,
      Function2<T,Out,java.util.concurrent.CompletionStage<T>> f);
Flow<In,java.util.List<Out>,Mat> groupedWithin(int n,
      scala.concurrent.duration.FiniteDuration d);
Flow<In,Out,Mat> log(java.lang.String name);
<T> Flow<In,T,Mat> map(Function<Out,T> f);
```

Using akka.streams.javadsl.Sink

```
static <T> Sink<T,java.util.concurrent.CompletionStage<Done>>
      foreach(Procedure<T> f);
static <U,In> Sink<In,java.util.concurrent.CompletionStage<U>>
      fold(U zero, Function2<U,In,U> f);
static <In> Sink<In,java.util.concurrent.CompletionStage<In>> head();
static <T> Sink<T,java.util.concurrent.CompletionStage<Done>> ignore();
static <In> Sink<In,java.util.concurrent.CompletionStage<In>> last();
// etc.
```

Example app - take two

Alpakka

Alpakka: Akka Streams connectors

- http://developer.lightbend.com/docs/alpakka/current/
- Example connectors: AMQP, AWS DynamoDB, AWS Kinesis, AWS Lambda, Cassandra, JMS, SSE, File IO, Azure, Camel, Kafka, TCP, etc.

Akka Actors & Streams

Akka Actors and Akka Streams

- Why actors?
  - Managing state.
- Why Akka Streams?
  - Process handling of "flowing" data.

Side note: Akka Streams uses Akka Actors under the hood, i.e. it is possible to implement "anything" in actors but it is more low level.

Stream -> Actor: *ask* with *mapAsync*

```
public static class MultiplierActor extends AbstractActor {
    @Override
    public Receive createReceive() {
        return receiveBuilder().match(Integer.class, i -> {
            getSender().tell(i * 2, getSelf());
        }).build();
    }
}
```

Stream -> Actor: *ask* with *mapAsync*

```
final ActorSystem actorSystem = ActorSystem.create();
final Materializer materializer =
    ActorMaterializer.create(actorSystem);
final ActorRef multiplier =
    actorSystem.actorOf(MultiplierActor.props());
Source.range(0, 10)
        .mapAsync(1, x -> ask(multiplier, x, 1000L))
        .map(e -> (Integer) e)
        .runWith(
             Sink.foreach(i ->
                  System.out.println("Example 3 - Number: " +
              i)), materializer);
```

Credits

- Colin Breck's inspirational blog post: http://blog.colinbreck.com/akka-streams-a-motivating-example
- My colleagues Björn Antonsson (@bantonsson), Peter Vlugter, Johan Andrén (@apnylle), and Konrad Malawski (@ktoso) for their help with this presentation/code.
- Lightbend for supporting this mission!


Lightbend